

lab03_partA

September 16, 2024

```
[25]: import pandas as pd
```

```
[26]: housing = pd.read_csv('housing.csv')
housing.head()
```

```
[26]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value
0	322.0	126.0	8.3252	452600.0
1	2401.0	1138.0	8.3014	358500.0
2	496.0	177.0	7.2574	352100.0
3	558.0	219.0	5.6431	341300.0
4	565.0	259.0	3.8462	342200.0

```
[27]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```

```
[28]: housing.describe()
```

```
[28]:
```

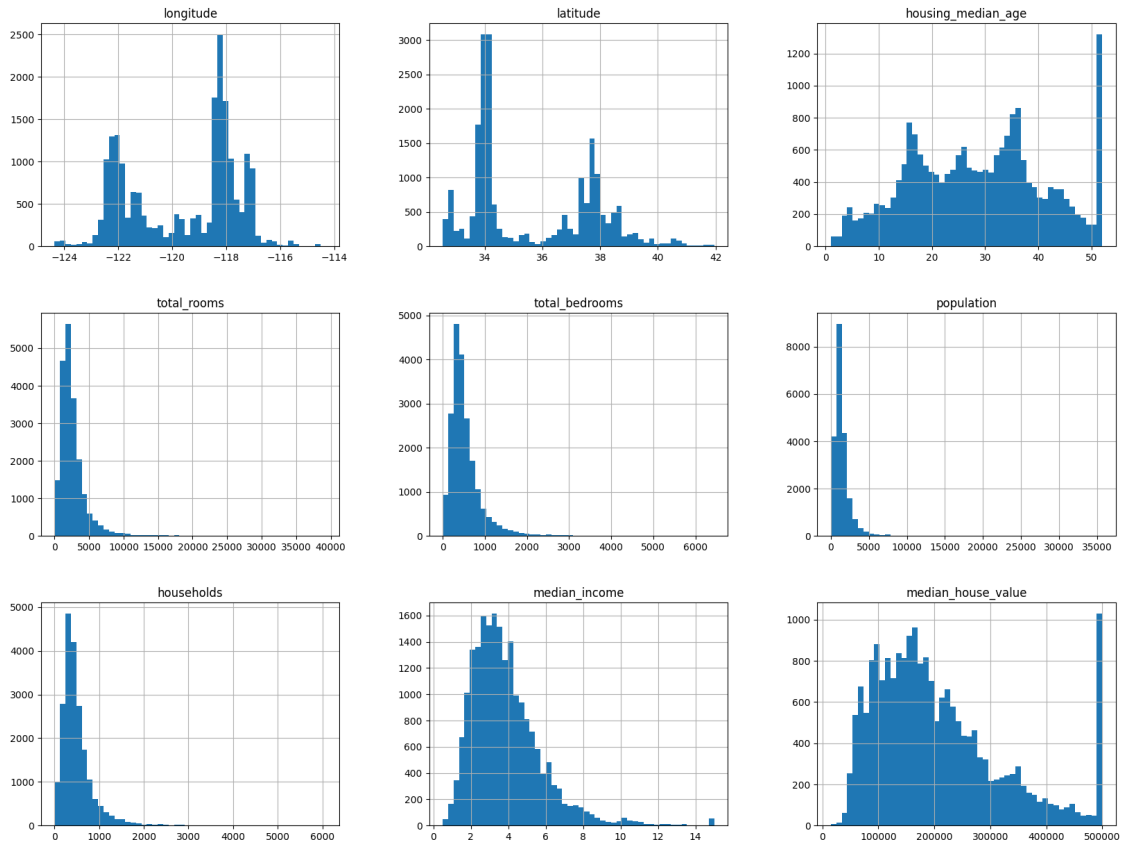
	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

0.1 Question 3

```
[29]: %matplotlib inline
import matplotlib.pyplot as plt
housing.hist(bins=50, figsize=(20,15))
plt.show()
```



0.2 Question 4

```
[30]: # to make this notebook's output identical and reproducible at every run, we
      ↪ need set seed
import numpy as np
np.random.seed(10)
```

```
[31]: # For illustration only. Sklearn has train_test_split()
def split_train_test(data, test_ratio):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size = int(len(data) * test_ratio)
    test_indices = shuffled_indices[:test_set_size]
    train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices], data.iloc[test_indices]
```

```
[32]: # run the function to get the train & test set
train_set, test_set = split_train_test(housing, 0.2)
```

```
[33]: train_samples = train_set.shape[0]
      test_samples = test_set.shape[0]
```

```

print(f"Number of samples in train_set: {train_samples}")
print(f"Number of samples in test_set: {test_samples}")
print(f"Total samples: {train_samples + test_samples}")

total_samples = housing.shape[0]
print(f"Total samples in the housing dataset: {total_samples}")

assert train_samples + test_samples == total_samples, "The total samples do not
↳match!"

```

```

Number of samples in train_set: 16512
Number of samples in test_set: 4128
Total samples: 20640
Total samples in the housing dataset: 20640

```

[34]: %pip install sklearn

```

Collecting sklearn
  Using cached sklearn-0.0.post12.tar.gz (2.6 kB)
  Preparing metadata (setup.py): started
  Preparing metadata (setup.py): finished with status 'error'
Note: you may need to restart the kernel to use updated packages.

error: subprocess-exited-with-error

× python setup.py egg_info did not run successfully.
  exit code: 1
  > [15 lines of output]
    The 'sklearn' PyPI package is deprecated, use 'scikit-learn'
    rather than 'sklearn' for pip commands.

Here is how to fix this error in the main use cases:
- use 'pip install scikit-learn' rather than 'pip install sklearn'
- replace 'sklearn' by 'scikit-learn' in your pip requirements files
  (requirements.txt, setup.py, setup.cfg, Pipfile, etc ...)
- if the 'sklearn' package is used by one of your dependencies,
  it would be great if you take some time to track which package uses
  'sklearn' instead of 'scikit-learn' and report it to their issue tracker
- as a last resort, set the environment variable
  SKLEARN_ALLOW_DEPRECATED_SKLEARN_PACKAGE_INSTALL=True to avoid this
error

More information is available at
https://github.com/scikit-learn/sklearn-pypi-package
[end of output]

```

note: This error originates from a subprocess, and is likely not a problem with pip.

error: metadata-generation-failed

× Encountered error while generating package metadata.
> See above for output.

note: This is an issue with the package mentioned above, not pip.

hint: See above for details.

[notice] A new release of pip is available: 24.1.2 -> 24.2

[notice] To update, run: python.exe -m pip install --upgrade pip

```
[35]: from sklearn.model_selection import train_test_split
      train_set, test_set = train_test_split(housing, test_size=0.2, random_state=10)
```

```
[36]: train_set[train_set.isna().any(axis=1)]
```

```
[36]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
6814	-118.07	34.10	32.0	4275.0	NaN	
4738	-118.38	34.05	49.0	702.0	NaN	
290	-122.16	37.77	47.0	1256.0	NaN	
19833	-119.38	36.53	38.0	1281.0	NaN	
4852	-118.31	34.03	47.0	1315.0	NaN	
...	
11512	-118.10	33.74	32.0	2035.0	NaN	
13656	-117.30	34.05	6.0	2155.0	NaN	
13015	-121.19	38.71	11.0	4415.0	NaN	
10236	-117.92	33.87	33.0	1597.0	NaN	
14521	-117.14	32.90	16.0	3217.0	NaN	

	population	households	median_income	median_house_value
6814	2812.0	1012.0	3.3512	214100.0
4738	458.0	187.0	4.8958	333600.0
290	570.0	218.0	4.3750	161900.0
19833	1423.0	293.0	1.9602	51400.0
4852	785.0	245.0	1.2300	138400.0
...
11512	934.0	512.0	4.2287	500001.0
13656	1039.0	391.0	1.6675	95800.0
13015	1520.0	627.0	3.2321	390800.0
10236	1888.0	423.0	3.0550	157800.0
14521	2054.0	687.0	4.2234	162100.0

[166 rows x 9 columns]

```
[37]: train_set_clean = train_set.dropna(subset=["total_bedrooms"])
      train_set_clean
```

```
[37]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
12346	-116.52	33.82	21.0	10227.0	2315.0	
19326	-122.94	38.50	46.0	2280.0	492.0	
16824	-122.49	37.63	31.0	3109.0	621.0	
6869	-118.11	34.06	16.0	2416.0	565.0	
16677	-120.66	35.13	41.0	2666.0	751.0	
...	
9372	-122.53	37.93	37.0	1722.0	352.0	
7291	-118.22	33.98	18.0	1781.0	765.0	
17728	-121.79	37.32	6.0	2850.0	561.0	
7293	-118.23	33.98	35.0	1366.0	496.0	
17673	-121.88	37.30	16.0	2692.0	749.0	

	population	households	median_income	median_house_value
12346	3623.0	1734.0	2.5212	145200.0
19326	807.0	366.0	2.6316	117000.0
16824	1472.0	618.0	5.1550	263900.0
6869	1750.0	514.0	2.8229	163700.0
16677	940.0	507.0	1.9653	236100.0
...
9372	648.0	337.0	4.1250	310300.0
7291	1913.0	702.0	1.2059	255000.0
17728	2160.0	581.0	5.5336	241900.0
7293	2160.0	497.0	2.2059	150000.0
17673	1674.0	681.0	2.6763	191100.0

[16346 rows x 9 columns]

```
[38]: train_labels = train_set_clean["median_house_value"].copy() # get labels for
      ↪ output label Y
      train_features = train_set_clean.drop("median_house_value", axis=1) # drop
      ↪ labels to get features X for training set
```

0.3 Question 8

Task 1

```
[39]: train_features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 16346 entries, 12346 to 17673
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              16346 non-null  float64
1   latitude               16346 non-null  float64
2   housing_median_age     16346 non-null  float64
3   total_rooms            16346 non-null  float64
4   total_bedrooms        16346 non-null  float64
```

```

5   population          16346 non-null   float64
6   households          16346 non-null   float64
7   median_income       16346 non-null   float64
dtypes: float64(8)
memory usage: 1.1 MB

```

Task 2

```

[40]: print(f"Shape of train_features: {train_features.shape}")

print("Missing values in train_features:")
print(train_features.isnull().sum())

print("Statistics summary of train_features:")
print(train_features.describe())

```

Shape of train_features: (16346, 8)

Missing values in train_features:

```

longitude      0
latitude       0
housing_median_age  0
total_rooms    0
total_bedrooms 0
population     0
households     0
median_income  0
dtype: int64

```

Statistics summary of train_features:

	longitude	latitude	housing_median_age	total_rooms \
count	16346.000000	16346.000000	16346.000000	16346.000000
mean	-119.576731	35.639888	28.606265	2630.196745
std	2.006389	2.139648	12.632219	2158.731496
min	-124.350000	32.540000	1.000000	2.000000
25%	-121.800000	33.930000	18.000000	1451.000000
50%	-118.500000	34.260000	29.000000	2125.500000
75%	-118.010000	37.720000	37.000000	3137.000000
max	-114.310000	41.950000	52.000000	39320.000000


	total_bedrooms	population	households	median_income
count	16346.000000	16346.000000	16346.000000	16346.000000
mean	537.625413	1425.149456	498.992659	3.864135
std	418.511077	1136.460782	378.999712	1.892231
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	785.000000	279.000000	2.559275
50%	435.000000	1165.500000	409.000000	3.533800
75%	647.000000	1728.000000	605.000000	4.739900
max	6445.000000	35682.000000	6082.000000	15.000100

```
[41]: print(train_labels.head())

import matplotlib.pyplot as plt

train_labels.hist(bins=50, figsize=(10, 6))
plt.xlabel("Median House Value")
plt.ylabel("Frequency")
plt.title("Histogram of Training Labels (Median House Value)")
plt.show()
```

```
12346    145200.0
19326    117000.0
16824    263900.0
6869     163700.0
16677    236100.0
Name: median_house_value, dtype: float64
```



```
[42]: train_features.hist(bins=50, figsize=(20,15))
```

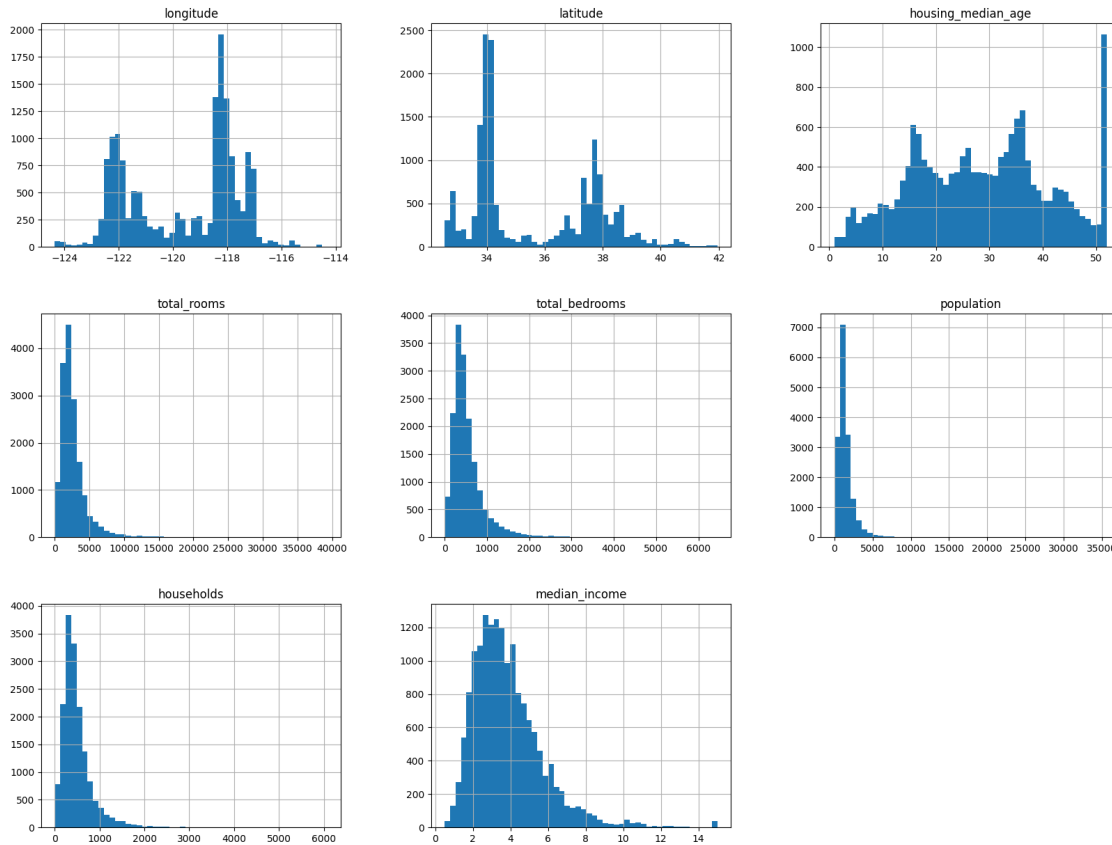
```
[42]: array([[<Axes: title={'center': 'longitude'}>,
<Axes: title={'center': 'latitude'}>,
<Axes: title={'center': 'housing_median_age'}>],
[<Axes: title={'center': 'total_rooms'}>,
<Axes: title={'center': 'total_bedrooms'}>],
```



```

<Axes: title={'center': 'population'}>],
[<Axes: title={'center': 'households'}>,
<Axes: title={'center': 'median_income'}>, <Axes: >]],
dtype=object)

```



```
[43]: train_features.describe()
```

```

[43]:      longitude      latitude  housing_median_age  total_rooms  \
count  16346.000000  16346.000000      16346.000000  16346.000000
mean    -119.576731    35.639888        28.606265    2630.196745
std       2.006389     2.139648        12.632219    2158.731496
min     -124.350000    32.540000         1.000000     2.000000
25%     -121.800000    33.930000        18.000000    1451.000000
50%     -118.500000    34.260000        29.000000    2125.500000
75%     -118.010000    37.720000        37.000000    3137.000000
max     -114.310000    41.950000        52.000000   39320.000000

      total_bedrooms  population  households  median_income
count  16346.000000  16346.000000  16346.000000  16346.000000
mean     537.625413   1425.149456    498.992659     3.864135

```

std	418.511077	1136.460782	378.999712	1.892231
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	785.000000	279.000000	2.559275
50%	435.000000	1165.500000	409.000000	3.533800
75%	647.000000	1728.000000	605.000000	4.739900
max	6445.000000	35682.000000	6082.000000	15.000100

```
[44]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler() ## define the transformer
scaler.fit(train_features) ## call .fit() method to calculate the min and max
      ↪ value for each column in datas
```

```
[44]: MinMaxScaler()
```

```
[45]: print("Min of each column: ",scaler.data_min_)
print("Max of each column: ",scaler.data_max_)
```

```
Min of each column:  [-124.35      32.54      1.          2.          1.          3.
1.
0.4999]
Max of each column:  [-1.14310e+02  4.19500e+01  5.20000e+01  3.93200e+04
6.44500e+03
3.56820e+04  6.08200e+03  1.50001e+01]
```

```
[46]: train_features_normalized = scaler.transform(train_features)
train_features_normalized
```

```
[46]: array([[0.77988048, 0.1360255 , 0.39215686, ..., 0.10146024, 0.28498602,
0.13939808],
[0.14043825, 0.63336876, 0.88235294, ..., 0.02253426, 0.06002302,
0.14701177],
[0.18525896, 0.54091392, 0.58823529, ..., 0.04117268, 0.10146358,
0.32103695],
...,
[0.25498008, 0.50797024, 0.09803922, ..., 0.06045573, 0.09537905,
0.34714694],
[0.60956175, 0.15302869, 0.66666667, ..., 0.06045573, 0.08156553,
0.11765355],
[0.24601594, 0.50584485, 0.29411765, ..., 0.04683427, 0.11182371,
0.15009448]])
```

0.4 Question 9

Task 1

```
[48]: # Assuming train_features_normalized is your normalized dataset

import matplotlib.pyplot as plt
import pandas as pd
```

```
# Convert the normalized NumPy array back to a DataFrame for plotting
normalized_df = pd.DataFrame(train_features_normalized, columns=train_features.
    ↪columns)

# Plot histograms for each normalized feature
normalized_df.hist(bins=50, figsize=(20, 15))
plt.show()
```

